

Amortized Analysis

Comp Sci 214, Fall 2022

The End is Near!

- Second midterm exam, Thu Dec 1st 11am, this room
 - ▶ Same format as first midterm, closed notes, no electronics
 - ▶ Not cumulative (but builds on earlier material)
 - ▶ Material until this Thursday is fair game
 - ▶ Will release a practice exam
- Final two lectures (Nov 22nd, 29th): cool, advanced topics!
 - ▶ Probabilistic data structures
 - ▶ Persistent data structures
- Final project
 - ▶ Initial deadline next Tuesday
 - ▶ Highly recommend getting (at least!) `locate_all` working
 - ▶ Keep a backup of your last non-crashing version!
 - ▶ Submit that to get actual feedback
 - ▶ Second try and report due Tue Dec 6th (finals week)

Previously on 214

When we saw union-finds (WQUPC), we never said how much a single union or find costs

Instead, we said that m operations on n objects is $\mathcal{O}(m \alpha(n))$

Previously on 214

When we saw union-finds (WQUPC), we never said how much a single union or find costs

Instead, we said that m operations on n objects is $\mathcal{O}(m \alpha(n))$

This is because some long-running operations do *maintenance* that make other operations *faster*

- I.e., find doing path compression

This means we can get a more accurate (and cheaper!) cost analysis by considering operations in *aggregate*, rather than assuming the worst case for each.

Intuition: with n operations, they can't all be the worst!

Today

- Two data structures where worst case is too pessimistic.
- If we consider cost of operations in bulk, we can get better complexity bounds!
→ *Amortized analysis*

Dynamic Arrays

Recall: dynamic arrays

Briefly mentioned all the way back in lecture 2.

- Data structure that implements the sequence ADT.
- Elements stored in a vector with spare capacity.
- When inserting, use that spare capacity.
- When we run out of capacity, allocate a bigger vector.

Back then, I said growing by doubling was efficient.

- Now let's see why!
- (Same reasoning applies to hash tables too.)

Recall: dynamic arrays

Briefly mentioned all the way back in lecture 2.

- Data structure that implements the sequence ADT.
- Elements stored in a vector with spare capacity.
- When inserting, use that spare capacity.
- When we run out of capacity, allocate a bigger vector.

Back then, I said growing by doubling was efficient.

- Now let's see why!
- (Same reasoning applies to hash tables too.)

But first, let's see why growing by just one element is inefficient.

A naïve implementation (1/2)

```
class BadDynArray[T]:  
  let data: VecC[T]
```

Store data in an array.

```
  def __init__(self):  
    self.data = []
```

Keep no spare capacity.

Most ops just call array ops.

```
  def len(self):  
    return self.data.len()
```

```
  def get(self, index):  
    return self.data[index]
```

```
  def set(self, index, element):  
    self.data[index] = element
```

```
  ...
```

A naïve implementation (2/2)

Grow by 1.

Create a new vector.

Copy elements over.

```
class BadDynArray[T]:
```

```
...
```

```
def push_back(self, element):  
    let new_data = [ None; self.len() + 1 ]  
    for i, v in self.data:  
        new_data[i] = v  
    new_data[self.len()] = element  
    self.data = new_data
```

Naïve representation complexities

- get/set/len are $\mathcal{O}(1)$
- push_back is $\mathcal{O}(n)!$

QUIZ: How long to build an n -element array, starting with an empty array and doing n pushes?

A: n

B: n^2

C: n^3

Naïve representation complexities

- get/set/len are $\mathcal{O}(1)$
- push_back is $\mathcal{O}(n)$!

QUIZ: How long to build an n -element array, starting with an empty array and doing n pushes?


A: n

B: n^2

C: n^3

$$\sum_{i=1}^n \mathcal{O}(i) = \mathcal{O}(n^2)$$

A better idea: leave extra space in the array



elements in dynamic array \neq vector length!

→ store size separately.

When we add an element, there may be spare capacity.

→ don't need to grow every time!

When we grow, grow enough to keep some spare capacity.

→ double when we need to grow.

Implementation (1/3)

```
class DynArray[T]:  
  let data: VecC[OrC(T, NoneC)] # could be unused space  
  let size: nat?  
  
  def __init__(self, initial_capacity: nat?):  
    self.data = [None; initial_capacity]  
    self.size = 0  
  
  def len(self):  
    return self.size  
  
  def capacity(self) -> nat?:  
    return self.data.len()
```

...

Not much to see here.

Implementation (2/3)

```
class DynArray[T]:  
  ...  
  
  def get(self, index):  
    self._bounds_check(index)  
    return self.data[index]  
  
  def set(self, index, element):  
    self._bounds_check(index)  
    self.data[index] = element  
  
  def _bounds_check(self, index):  
    if index >= self.size:  
      error('DynArray: out of bounds')  
  
  ...
```

Need to check that we're not accessing unused spots!

Implementation (3/3)

```
class DynArray[T]:  
  ...  
  
  def push_back(self, element):  
    self._ensure_capacity(self.size + 1)  
    self.data[self.size] = element  
    self.size = self.size + 1  
  
  def _ensure_capacity(self, cap):  
    if self.capacity() < cap:  
      cap = max(cap, 2 * self.capacity())  
      let new_data = [ None; cap ]  
      for i, v in self.data:  
        new_data[i] = v  
      self.data = new_data
```

Double when we need to grow.

Time complexities

- `get/set/len` are $\mathcal{O}(1)$
- `push_back` is still $\mathcal{O}(n)$ in the worst case!

QUIZ: How long to build an n -element array, starting with an empty array and doing n pushes?

A: n

B: n^2

C: n^3

D: It's complicated

Time complexities

- get/set/len are $\mathcal{O}(1)$
- push_back is still $\mathcal{O}(n)$ in the worst case!

QUIZ: How long to build an n -element array, starting with an empty array and doing n pushes?

A: n

B: n^2

C: n^3


D: It's complicated

$$\sum_{i=0}^n \mathcal{O}(i) = \mathcal{O}(n^2)?$$

That's if all pushes are the worst case. Will they be?

The peculiar thing about push_back


- Most of the time it's cheap
- But *occasionally*, we need to grow (which is expensive)



X axis: push #

Y axis: time it takes to do push #x

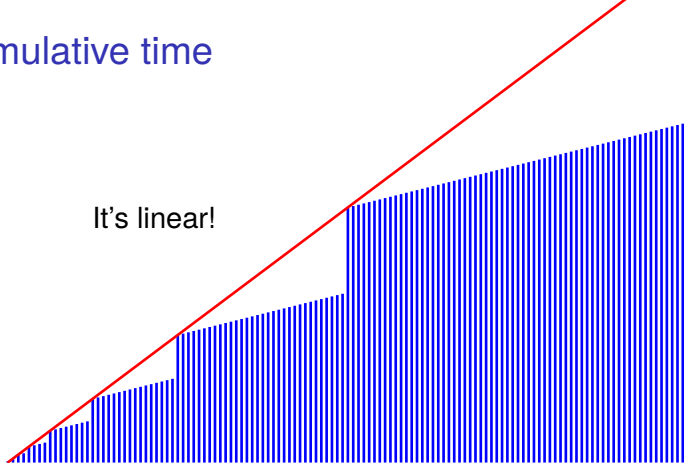
Cumulative time



X axis: push #

Y axis: *cumulative* time it takes to do pushes $0 \dots x$

Cumulative time



X axis: push #

Y axis: *cumulative* time it takes to do pushes $0 \dots x$

Wait, what?

- push_back has two “modes”
 - ▶ There is spare capacity: set array element, $\mathcal{O}(1)$
 - ▶ There is no spare capacity: double the size, $\mathcal{O}(n)$
 - ▶ As we go from 0 to n elements, mix of both modes
- When we double, get space for **many** cheap pushes!
 - ▶ If we currently have m elements, get m cheap pushes
 - ▶ Doublings get less and less frequent as array grows
- Now let's work out the math of this

Dynamic array aggregate analysis

Suppose we create a new array and push n times. How can we show linear total time?

Dynamic array aggregate analysis

Suppose we create a new array and push n times. How can we show linear total time?

Let s_i be the size of the array at step i .

Let c_i be the cost of the i th push:

$$c_i = \begin{cases} i & \text{if } i - 1 \text{ is a power of 2 (i.e., we're full)} \\ 1 & \text{otherwise (i.e., we have spare capacity)} \end{cases}$$

i	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17
s_i	1	2	4	4	8	8	8	8	16	16	16	16	16	16	16	16	32
c_i	1	2	3	1	5	1	1	1	9	1	1	1	1	1	1	1	17

Adding it up

Let $d_i = c_i - 1$ (the cost of doubling, when we do double)

Then,

$$\sum_{i=1}^n c_i = \sum_{i=1}^n (1 + d_i)$$

Adding it up

Let $d_i = c_i - 1$ (the cost of doubling, when we do double)

Then,

$$\begin{aligned}\sum_{i=1}^n c_i &= \sum_{i=1}^n (1 + d_i) \\ &= n + \sum_{i=1}^n d_i\end{aligned}$$

d_i is almost always 0.

Adding it up

Let $d_i = c_i - 1$ (the cost of doubling, when we do double)

Then,

$$\sum_{i=1}^n c_i = \sum_{i=1}^n (1 + d_i)$$

$$= n + \sum_{i=1}^n d_i$$

$$= n + \sum_{i=0}^{\log_2 n} 2^i$$

d_i is almost always 0.

We double $\log_2 n$ times.

Adding it up

Let $d_i = c_i - 1$ (the cost of doubling, when we do double)

Then,

$$\sum_{i=1}^n c_i = \sum_{i=1}^n (1 + d_i)$$

$$= n + \sum_{i=1}^n d_i$$

d_i is almost always 0.

$$= n + \sum_{i=0}^{\log_2 n} 2^i$$

We double $\log_2 n$ times.

$$= n + (1 + 2 + 4 + \dots + \frac{n}{4} + \frac{n}{2} + n)$$

Adding it up

Let $d_i = c_i - 1$ (the cost of doubling, when we do double)

Then,

$$\sum_{i=1}^n c_i = \sum_{i=1}^n (1 + d_i)$$

$$= n + \sum_{i=1}^n d_i$$

d_i is almost always 0.

$$= n + \sum_{i=0}^{\log_2 n} 2^i$$

We double $\log_2 n$ times.

$$= n + (1 + 2 + 4 + \dots + \frac{n}{4} + \frac{n}{2} + n) \\ < 3n$$

Time complexities

- get/set/size are $\mathcal{O}(1)$
- push_back has $\mathcal{O}(1)$ *amortized cost*
 - ▶ If you do n of them, the *total* cost is $\mathcal{O}(n)$

Stop! Question time!

Before we move on to our second data structure...

Anything unclear so far?

Anything I missed?

Aside: Do we *have* to double?

- Doubling each time is efficient
- Adding 1 is not
- What about other options?

Aside: Do we *have* to double?

- Doubling each time is efficient
- Adding 1 is not
- What about other options?
- Any *multiplicative* growth works!
 - ▶ Tripling each time
 - ▶ Growing by $1.5\times$
 - ▶ Even growing by just $1.0001\times$!
 - ▶ All amortized $\mathcal{O}(1)$; just different constant factors
 - ▶ In practice, $1.5\times$ is a common choice
- Adding a constant amount of space doesn't work
 - ▶ Not even adding 1,000,000 each time we grow!
 - ▶ That only gives us a *constant* number of cheap pushes to amortize a $\mathcal{O}(n)$ operation over!

Banker's Queues

Banker's queues

- Another data structure that implements the queue ADT
- Will have queue operations with $\mathcal{O}(1)$ amortized cost
- Unlike other queues we've seen, easy to make *persistent*
 - ▶ More on persistent data structures in a later lecture

Banker's queues

- Another data structure that implements the queue ADT
- Will have queue operations with $\mathcal{O}(1)$ amortized cost
- Unlike other queues we've seen, easy to make *persistent*
 - ▶ More on persistent data structures in a later lecture

Key idea: use two stacks, *front* and *back*

- Enqueue into *back*, dequeue from *front*
- When *front* is empty, transfer contents of *back* into it
- The elements of the queue go from top of front to bottom of front, then bottom of back to top of back

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $|3\rangle$,

back: $|4, 5\rangle$ }

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $|3\rangle$,

back: $|4, 5\rangle$ }

`q.dequeue()` \Rightarrow ?

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle \}$

`q.dequeue()` \Rightarrow 3

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle \}$

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle \}$

`q.dequeue()` $\Rightarrow 3$

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle \}$

`q.enqueue(6)`

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle \}$

q.dequeue() $\Rightarrow 3$

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle \}$

q.enqueue(6)

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle \}$

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle \}$

`q.dequeue()` \Rightarrow 3

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle \}$

`q.enqueue(6)`

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle \}$

`q.dequeue()` \Rightarrow ?

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle \}$

`q.dequeue()` $\Rightarrow 3$

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle \}$

`q.enqueue(6)`

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle \}$

`q.dequeue()` $\Rightarrow ?$

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle \}$

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle \}$

`q.dequeue()` $\Rightarrow 3$

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle \}$

`q.enqueue(6)`

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle \}$

`q.dequeue()` $\Rightarrow ?$

$\langle 4, 5, 6 \langle$

{front: $| 6 \rangle$,

back: $| 4, 5 \rangle \}$

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle \}$

q.dequeue() $\Rightarrow 3$

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle \}$

q.enqueue(6)

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle \}$

q.dequeue() $\Rightarrow ?$

$\langle 4, 5, 6 \langle$

{front: $| 6, 5 \rangle$,

back: $| 4 \rangle \}$

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle \}$

`q.dequeue()` $\Rightarrow 3$

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle \}$

`q.enqueue(6)`

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle \}$

`q.dequeue()` $\Rightarrow ?$

$\langle 4, 5, 6 \langle$

{front: $| 6, 5, 4 \rangle$,

back: $| \rangle \}$

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle \}$

q.dequeue() $\Rightarrow 3$

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle \}$

q.enqueue(6)

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle \}$

q.dequeue() $\Rightarrow 4$

$\langle 4, 5, 6 \langle$

{front: $| 6, 5, 4 \rangle$,

back: $| \rangle \}$

$\langle 5, 6 \langle$

{front: $| 6, 5 \rangle$,

back: $| \rangle \}$

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle$ }

q.dequeue() \Rightarrow 3

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle$ }

q.enqueue(6)

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle$ }

q.dequeue() \Rightarrow 4

$\langle 4, 5, 6 \langle$

{front: $| 6, 5, 4 \rangle$,

back: $| \rangle$ }

$\langle 5, 6 \langle$

{front: $| 6, 5 \rangle$,

back: $| \rangle$ }

q.enqueue(7)

Banker's queue example

$\langle 3, 4, 5 \langle$

{front: $| 3 \rangle$,

back: $| 4, 5 \rangle$ }

q.dequeue() $\Rightarrow 3$

$\langle 4, 5 \langle$

{front: $| \rangle$,

back: $| 4, 5 \rangle$ }

q.enqueue(6)

$\langle 4, 5, 6 \langle$

{front: $| \rangle$,

back: $| 4, 5, 6 \rangle$ }

q.dequeue() $\Rightarrow 4$

$\langle 4, 5, 6 \langle$

{front: $| 6, 5, 4 \rangle$,

back: $| \rangle$ }

$\langle 5, 6 \langle$

{front: $| 6, 5 \rangle$,

back: $| \rangle$ }

q.enqueue(7)

$\langle 5, 6, 7 \langle$

{front: $| 6, 5 \rangle$,

back: $| 7 \rangle$ }

Banker's queue implementation (1/2)

```
class BankersQueue[T] (QUEUE):  
  let front  
  let back  
  # Interpretation: the queue is the elements of  
  # `front` in pop order followed by `back` in reverse  
  
  def __init__(self):  
    self.front = Stack()  
    self.back = Stack()  
  
  def len(self):  
    return self.front.len() + self.back.len()  
  
  def empty?(self):  
    return self.front.empty?() and self.back.empty?()  
  
  ...
```

Banker's queue implementation (2/2)

```
class BankersQueue[T] (QUEUE):  
  ...  
  
  def enqueue(self, element: T) -> NoneC:  
    self.back.push(element)
```

Banker's queue implementation (2/2)

```
class BankersQueue[T] (QUEUE):  
    ...  
  
    def enqueue(self, element: T) -> NoneC:  
        self.back.push(element)  
  
    def dequeue(self) -> T:  
        # refill `front` if needed  
        if self.front.empty():  
            if self.back.empty():  
                error('BankersQueue.dequeue: empty')  
            while not self.back.empty():  
                # reverses order! top of `back` will  
                # end up on bottom of `front`  
                self.front.push(self.back.pop())  
        return self.front.pop()
```

We have both fast case and slow case → amortized analysis!

Stop! Question time!

Before we move on to the analysis...

Anything unclear so far?

Anything I missed?

Banker's queue analysis

We assign a “balance” (as in, bank account balance) to each data structure state:

$$B(q) = q.\text{back}.\text{len}()$$

Operations may deposit to or withdraw from the balance.

Balance of a new queue is 0, and balance is never negative.

Banker's queue analysis

We assign a “balance” (as in, bank account balance) to each data structure state:

$$B(q) = q.\text{back}.\text{len}()$$

Operations may deposit to or withdraw from the balance.

Balance of a new queue is 0, and balance is never negative.

The amortized cost of an operation will be

$$c + B(q') - B(q)$$

where c is the actual cost of the operation,

q is the state before the operation, and q' is after.

Balance change is $+$ for deposits, $-$ for withdrawals

Key idea: some operations will spend savings to get an amortized cost lower than their actual cost.

Actual costs

Actual cost of enqueue: 1

Actual costs

Actual cost of enqueue: 1

Actual cost of cheap dequeue (front isn't empty): 1

Actual costs

Actual cost of enqueue: 1

Actual cost of cheap dequeue (front isn't empty): 1

Actual cost of expensive dequeue (with reversal):

- Cost of the reversal (the number of elements reversed): n
- Plus the cost of a cheap dequeue: 1
- Total: $n + 1$

Amortized cost of enqueue

- Actual cost of enqueue is 1
- Increases the length of the back by 1, hence
$$B(q') - B(q) = 1$$
 - ▶ I.e., *deposit* of 1

So amortized cost is $1 + 1 = 2$. That's $\mathcal{O}(1)$.

Amortized cost of cheap dequeue

- Actual cost of cheap dequeue is 1
- No change in balance; no change to back

So amortized cost is 1. Also $\mathcal{O}(1)$.

Amortized cost of expensive dequeue

Let n be $q.\text{back}.\text{len}()$, the length of the back stack. Then:

- Actual cost is $n + 1$
- $B(q) = n$ (before reversal)
- $B(q') = 0$ (after reversal)
- *Withdrawal* of n

So amortized cost is $(n + 1) + (0 - n) = 1$.

Look at that! $\mathcal{O}(1)$!

Wait, what just happened?

operation	single operation*	amortized
enqueue	$\mathcal{O}(1)$	$\mathcal{O}(1)$
dequeue	$\mathcal{O}(n)$	$\mathcal{O}(1)$

We found a way (B) to account for cost across operations.

- No fraud, though! All costs are accounted for somewhere.
- We're just shifting costs from expensive ops to cheap ones.

* worst-case

Wait, what just happened?

operation	single operation*	amortized
enqueue	$\mathcal{O}(1)$	$\mathcal{O}(1)$
dequeue	$\mathcal{O}(n)$	$\mathcal{O}(1)$

We found a way (B) to account for cost across operations.

- No fraud, though! All costs are accounted for somewhere.
- We're just shifting costs from expensive ops to cheap ones.

Insight: dequeue can only be expensive if we did a lot of (cheap) enqueues before!

We're "saving up" time on cheap operations, and "spending" it on the expensive ones. So everything works out constant.

* worst-case

Amortized data structures

data structure	operation	single operation*	amortized
dynamic array	push_back	$\mathcal{O}(n)$	$\mathcal{O}(1)$
banker's queue	enqueue	$\mathcal{O}(1)$	$\mathcal{O}(1)$
	dequeue	$\mathcal{O}(n)$	$\mathcal{O}(1)$
WQUPC	union	$\mathcal{O}(\log n)$	$\mathcal{O}(\alpha(n))^{**}$
	find	$\mathcal{O}(\log n)$	$\mathcal{O}(\alpha(n))^{**}$

* worst-case

** AKA “I can’t believe it’s not constant!”

If we have time: playing with our dynamic array
and/or banker's queue

Next time: The relational model